

Course 1

# Introduction to Machine Learning

Christophe Eloy

# Outline

- Practical information
- What is machine learning?
- Different types of learning
- Examples
- Linear models
  - Linear regression with one variable
  - Linear regression with multiple variables

# Practical information

- Alternance of classes and tutorials
- Tutorials with Python/Jupyter + standard libraries (numpy, pandas, etc.)
- Assessment with a final presentation/project
- Contact me: [christophe.elay@univ-amu.fr](mailto:christophe.elay@univ-amu.fr)

# Useful resources

- MOOC Machine Learning, Andrew Ng (Stanford)
- Review paper “A high-bias, low-variance introduction to Machine Learning for physicists”, Pankaj Mehta et al. with Jupyter notebooks
- Lectures on Learning from Data, Yaser Abu-Mostafa (Caltech) on iTunesU and YouTube

# What is Machine Learning?

- Two definitions


**“The field of study that gives computers the ability to learn without being explicitly programmed.”**

**—Arthur Samuel**

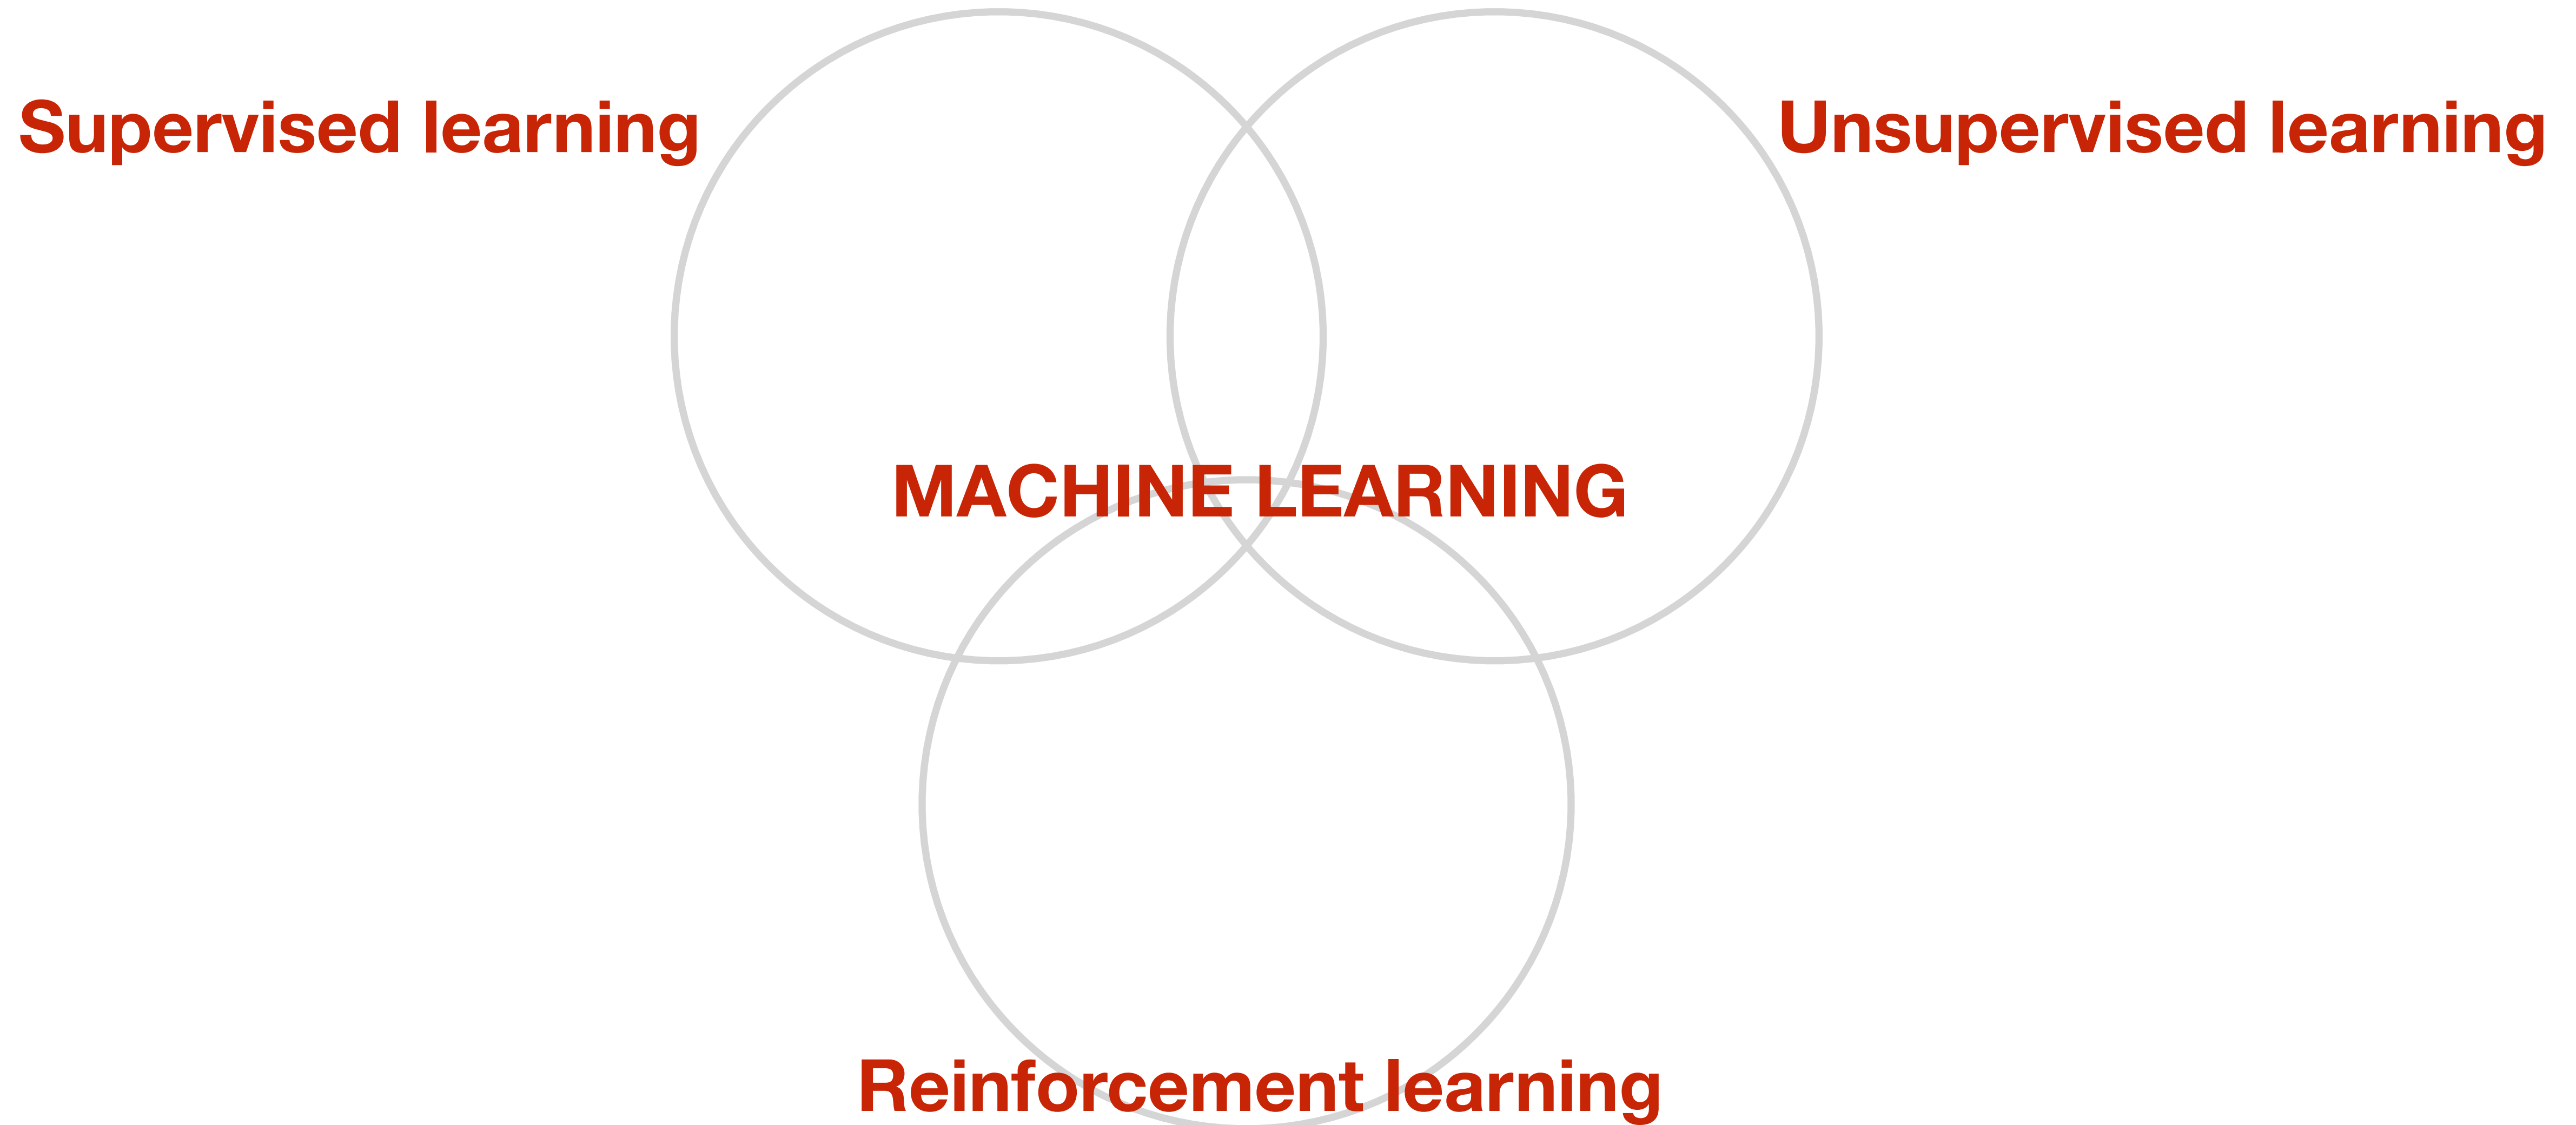
**"A computer program is said to learn from experience  $E$  with respect to some class of tasks  $T$  and performance measure  $P$ , if its performance at tasks in  $T$ , as measured by  $P$ , improves with experience  $E$ ."**

**—Tom Mitchell**

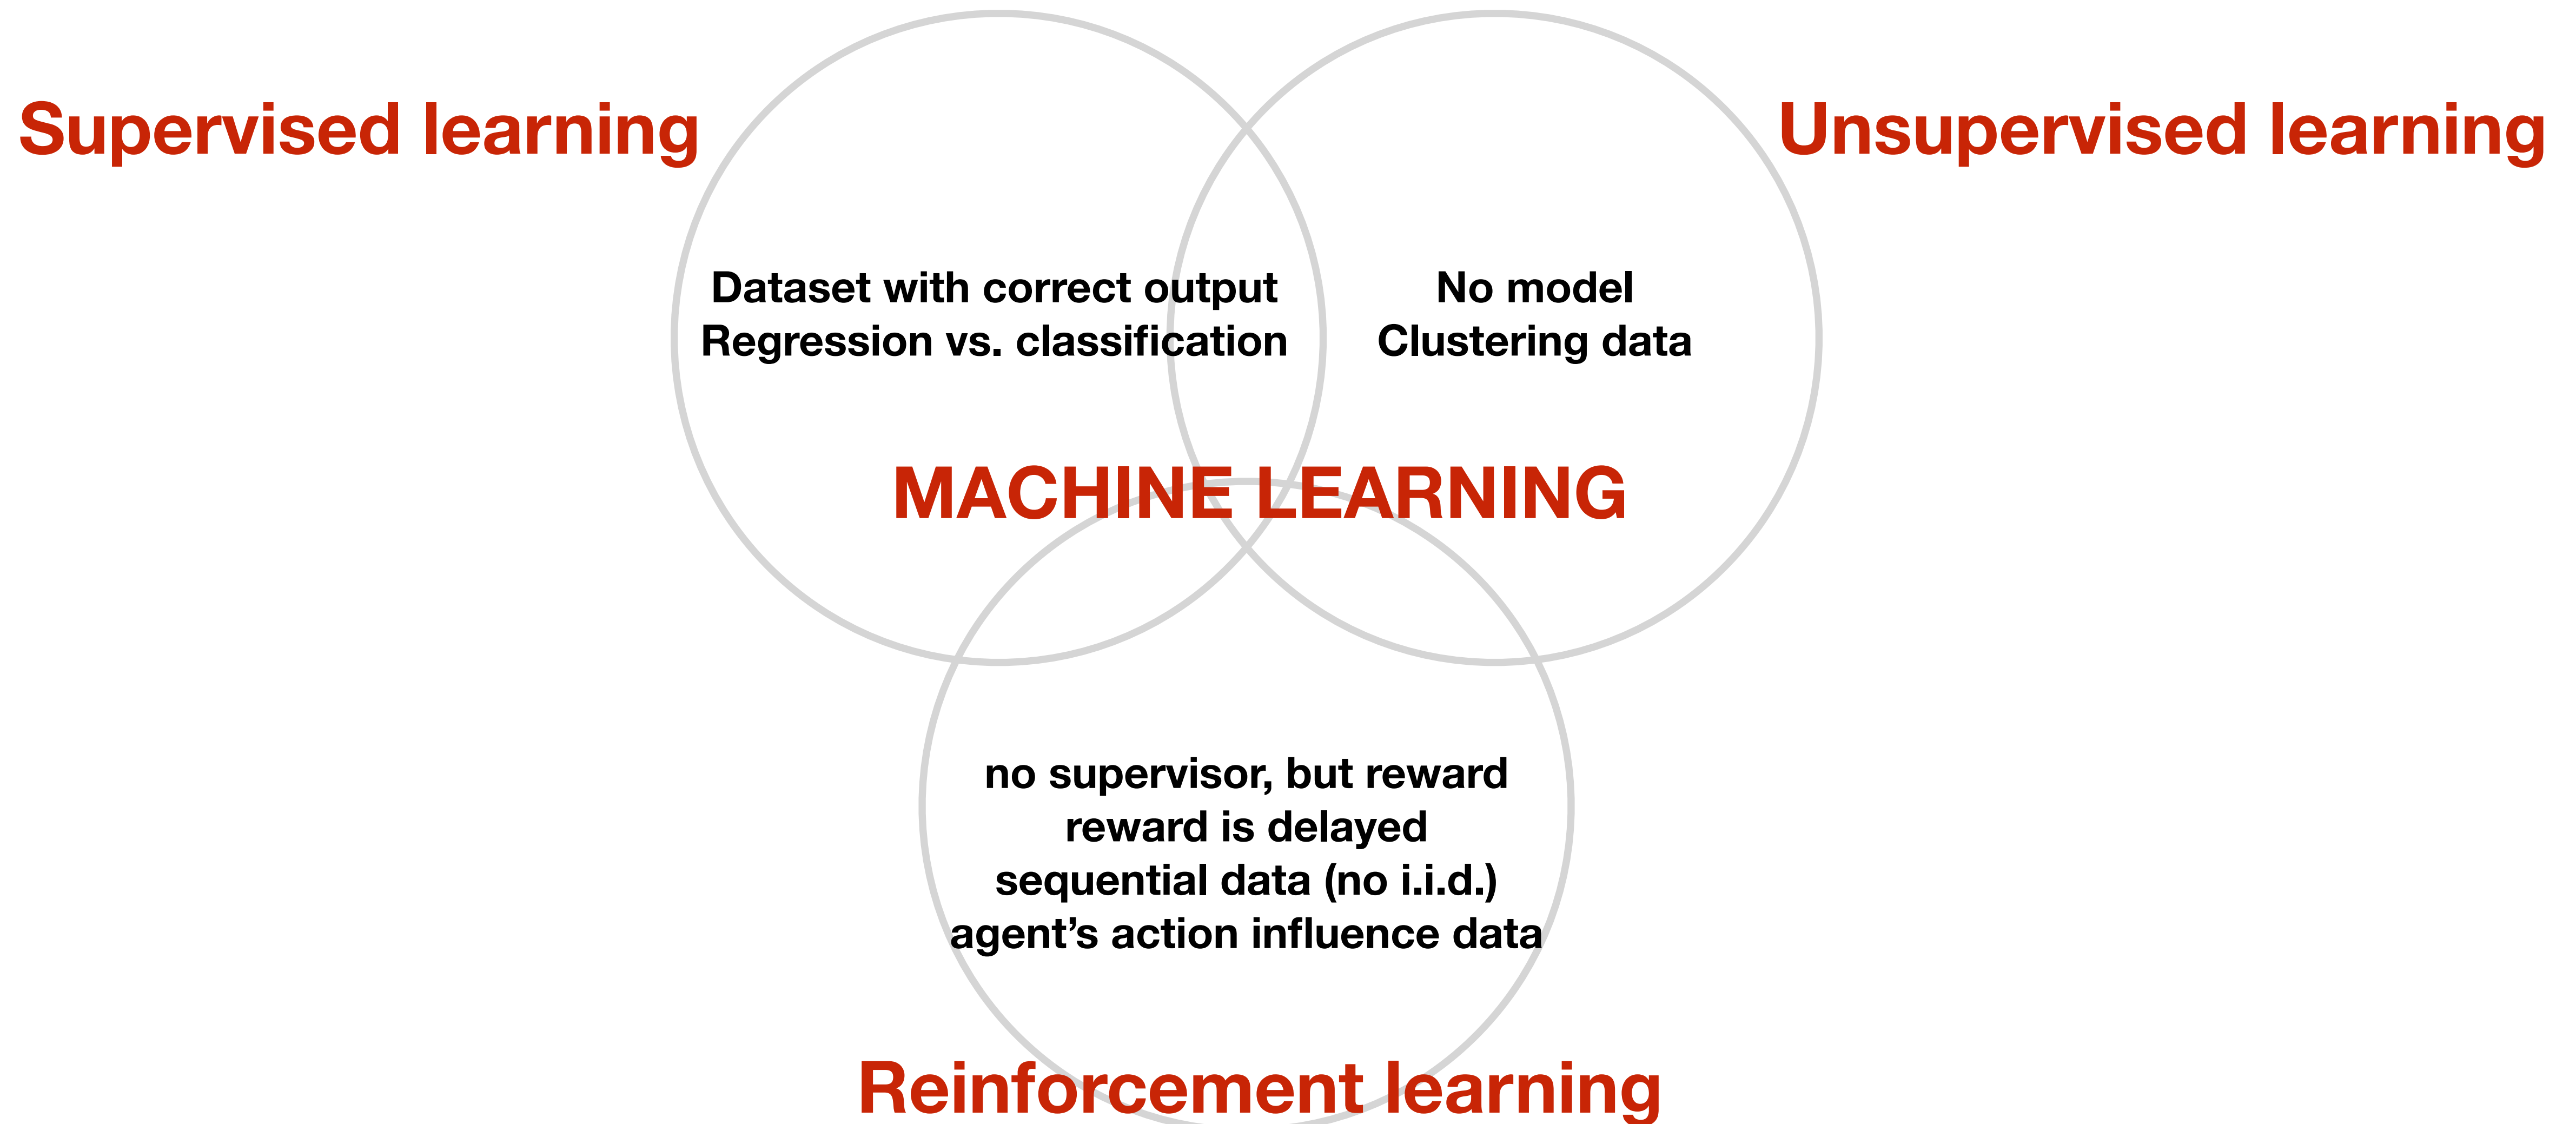
# At the crossroad of many fields



# Branches of Machine Learning



# Branches of Machine Learning





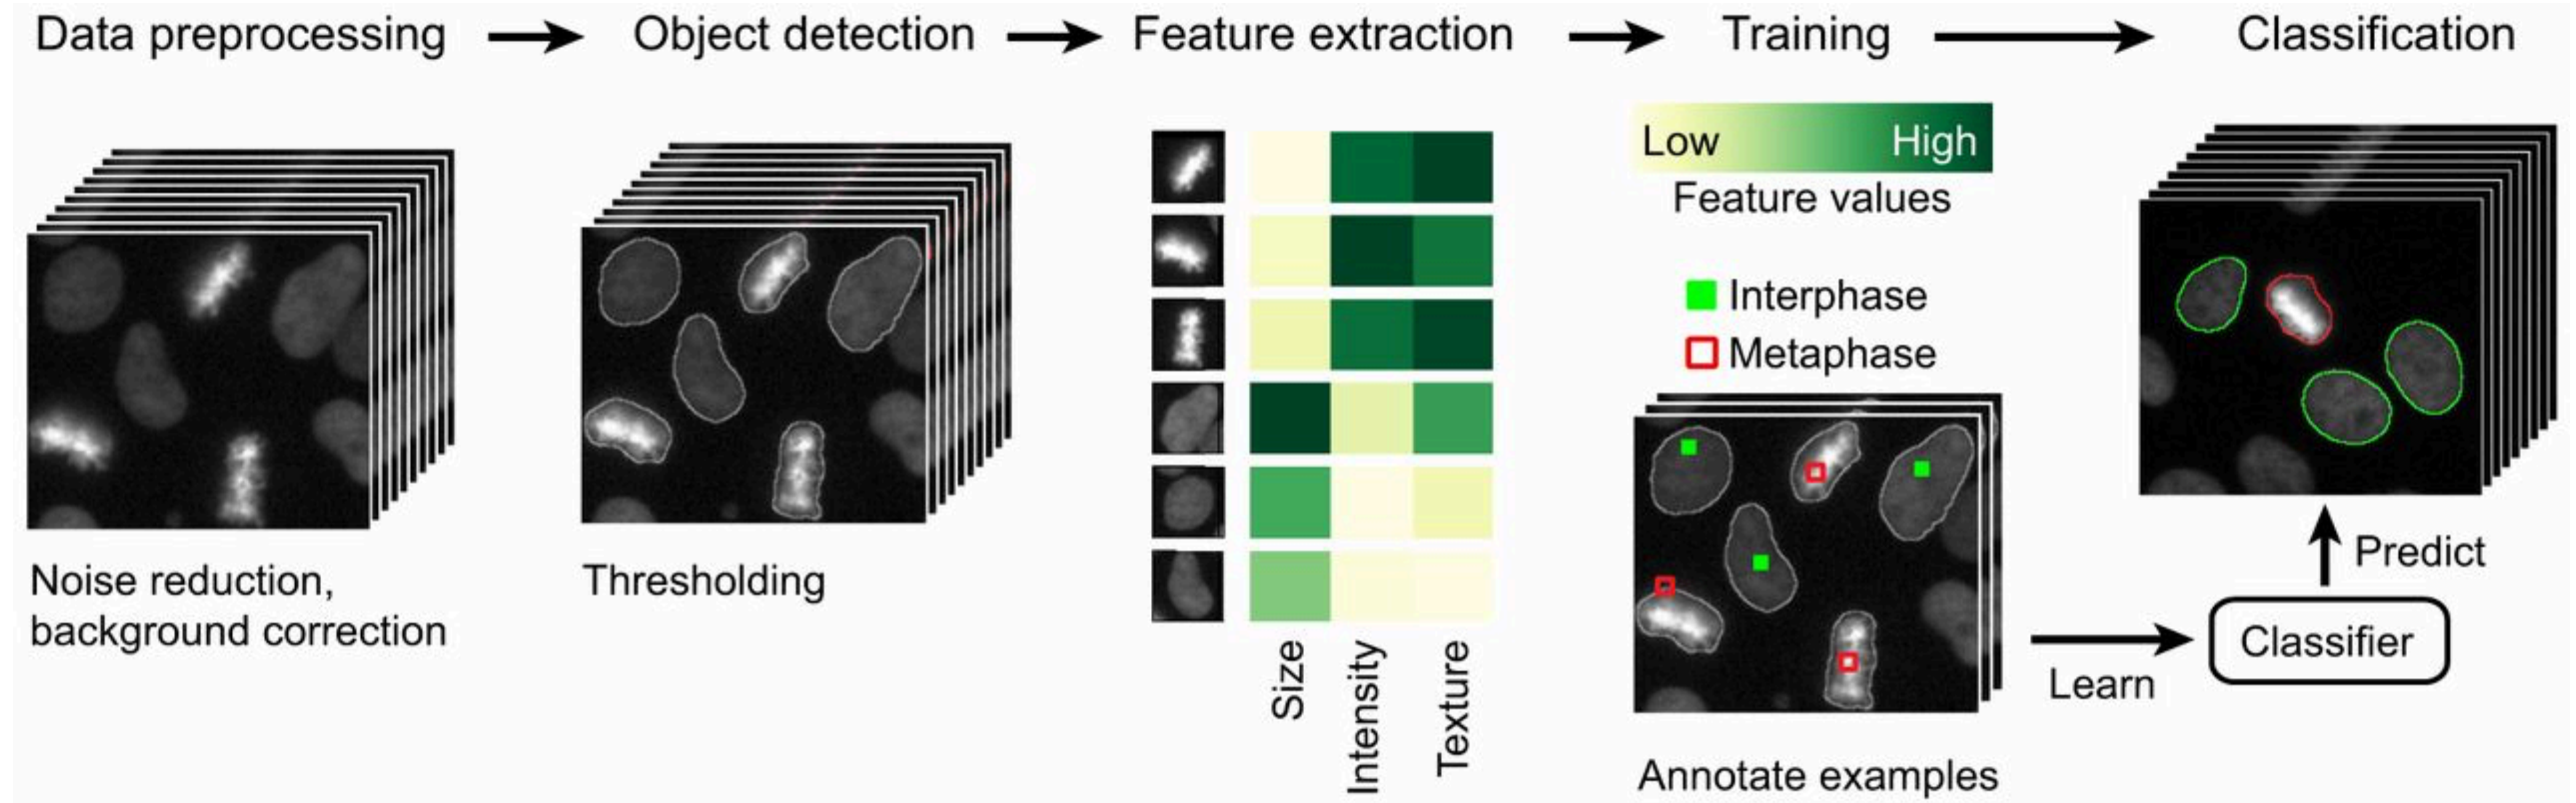
# Examples

- Classify emails as spam or not
- Predict the success of a movie
- Diagnose from a list of symptoms
- Drive a car autonomously
- Translate a text
- Find groups in a social network
- Defeat the world champion of Go

# Examples


- Classify emails as spam or not      Classification (supervised learning)
- Predict the success of a movie      Regression (supervised learning)
- Diagnose from a list of symptoms      Classification (supervised learning)
- Drive a car autonomously      Reinforcement learning
- Translate a text      Reinforcement learning
- Find groups in a social network      Unsupervised learning
- Defeat the world champion of Go      Reinforcement learning

# Recognize and classify phenotypes

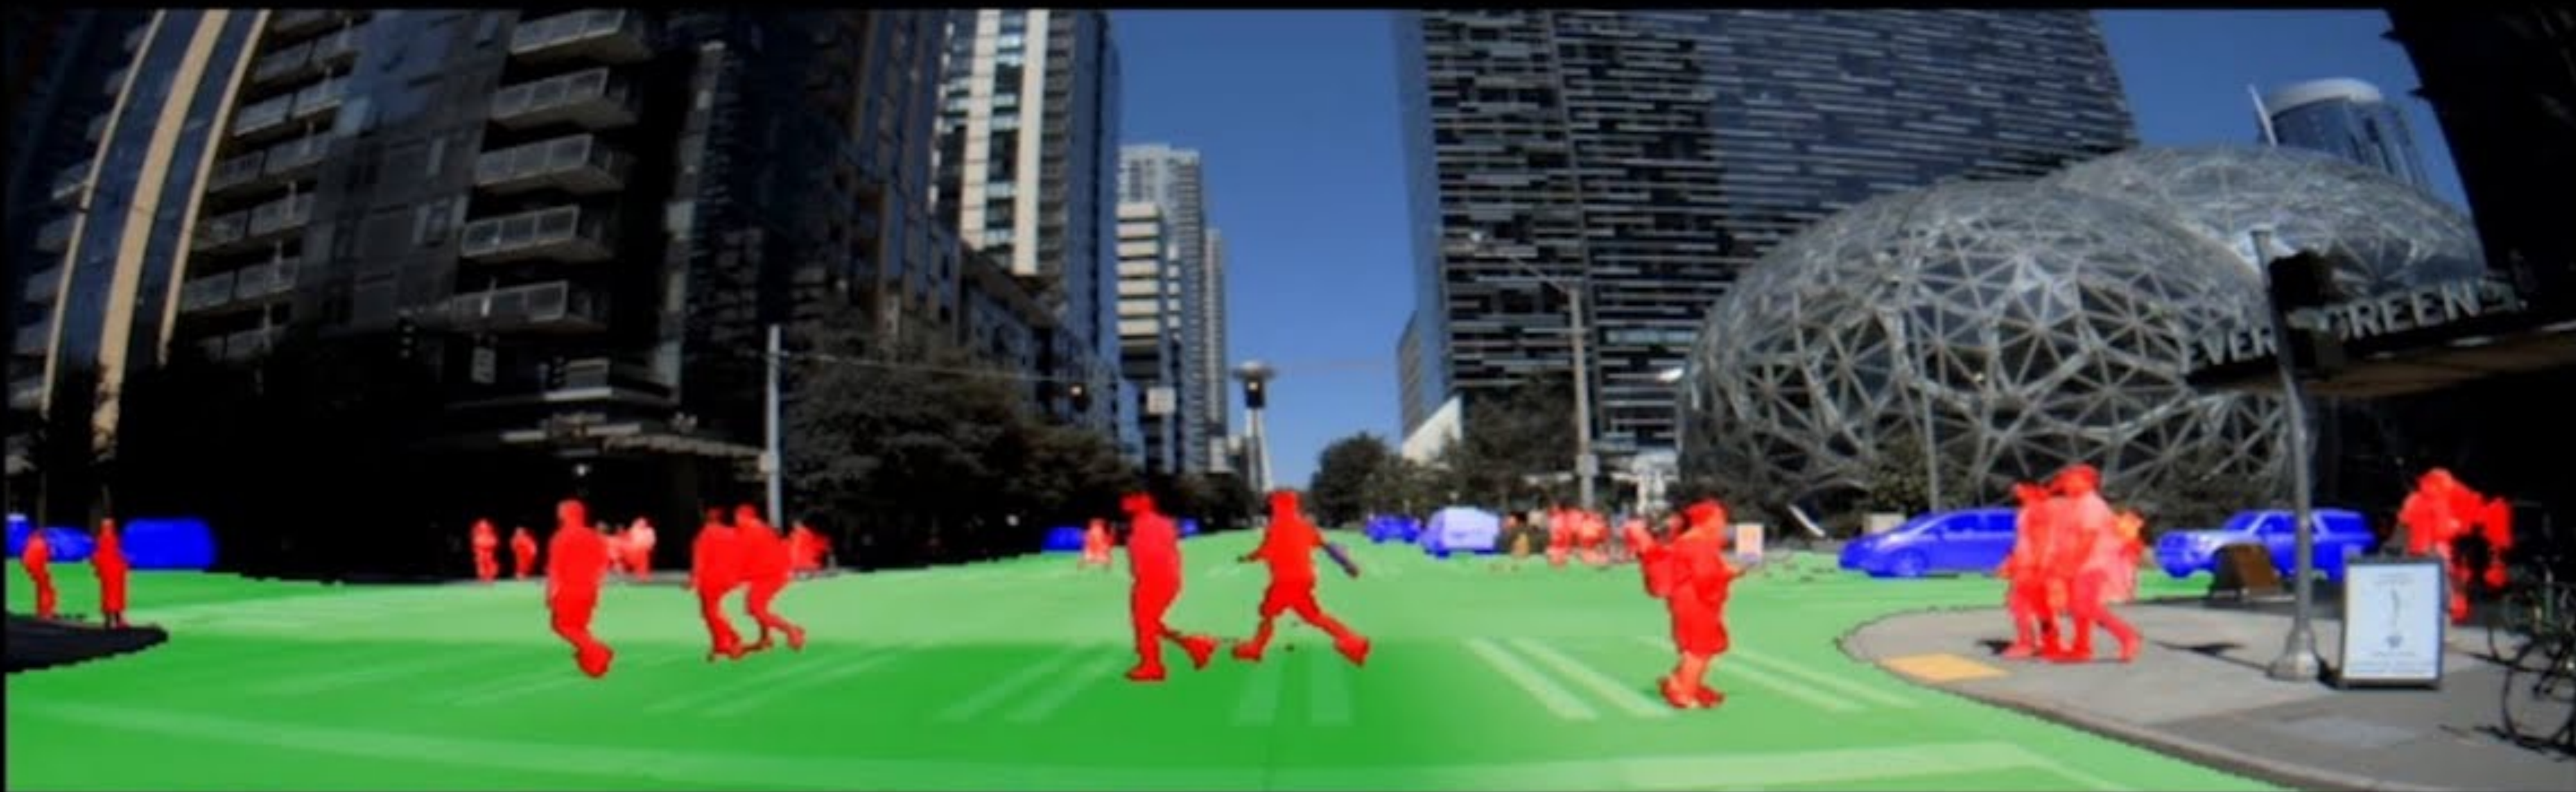




# Learning to exploit vortices











ima...

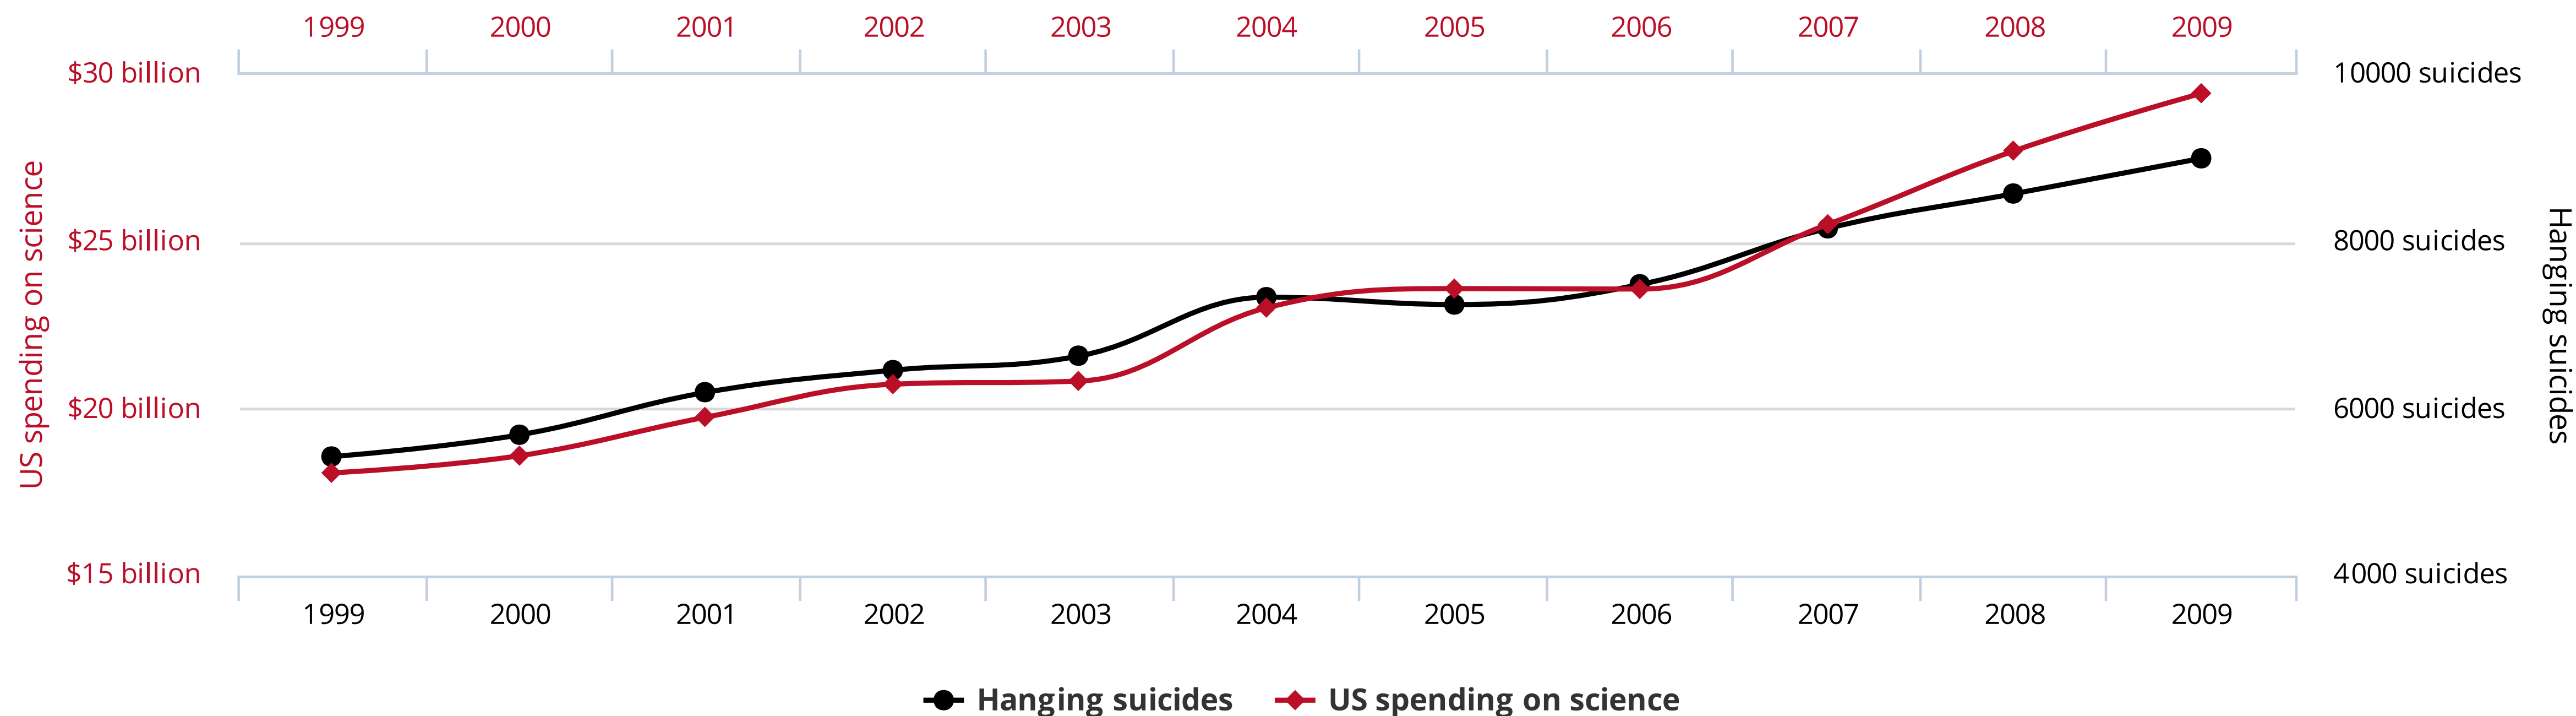


021 3 1

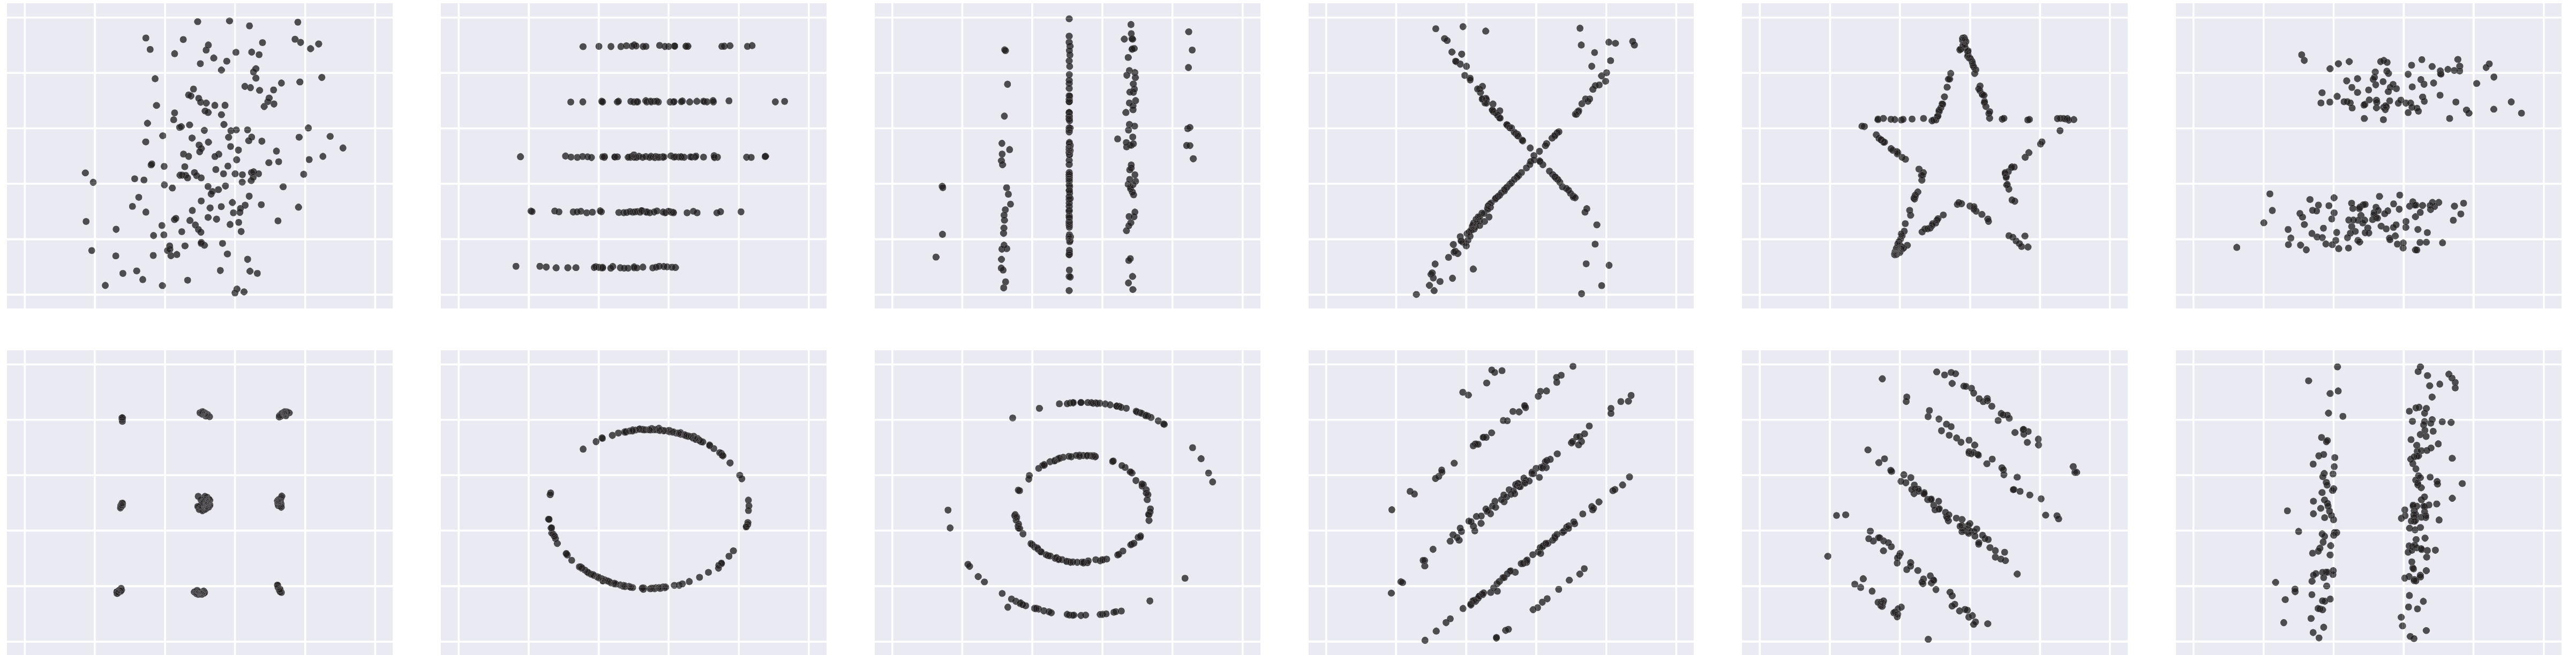


# Beware of spurious correlations!

**US spending on science, space, and technology**  
correlates with  
**Suicides by hanging, strangulation and suffocation**



# Plot your data!



**All data have same means, same std, same Pearson's correlation coefficient  $r$**




# When to use Machine Learning?

- There is a “pattern”
- Relationships are not tractable mathematically
- There is (a lot of) data

# Univariate linear regression



# Simple working example

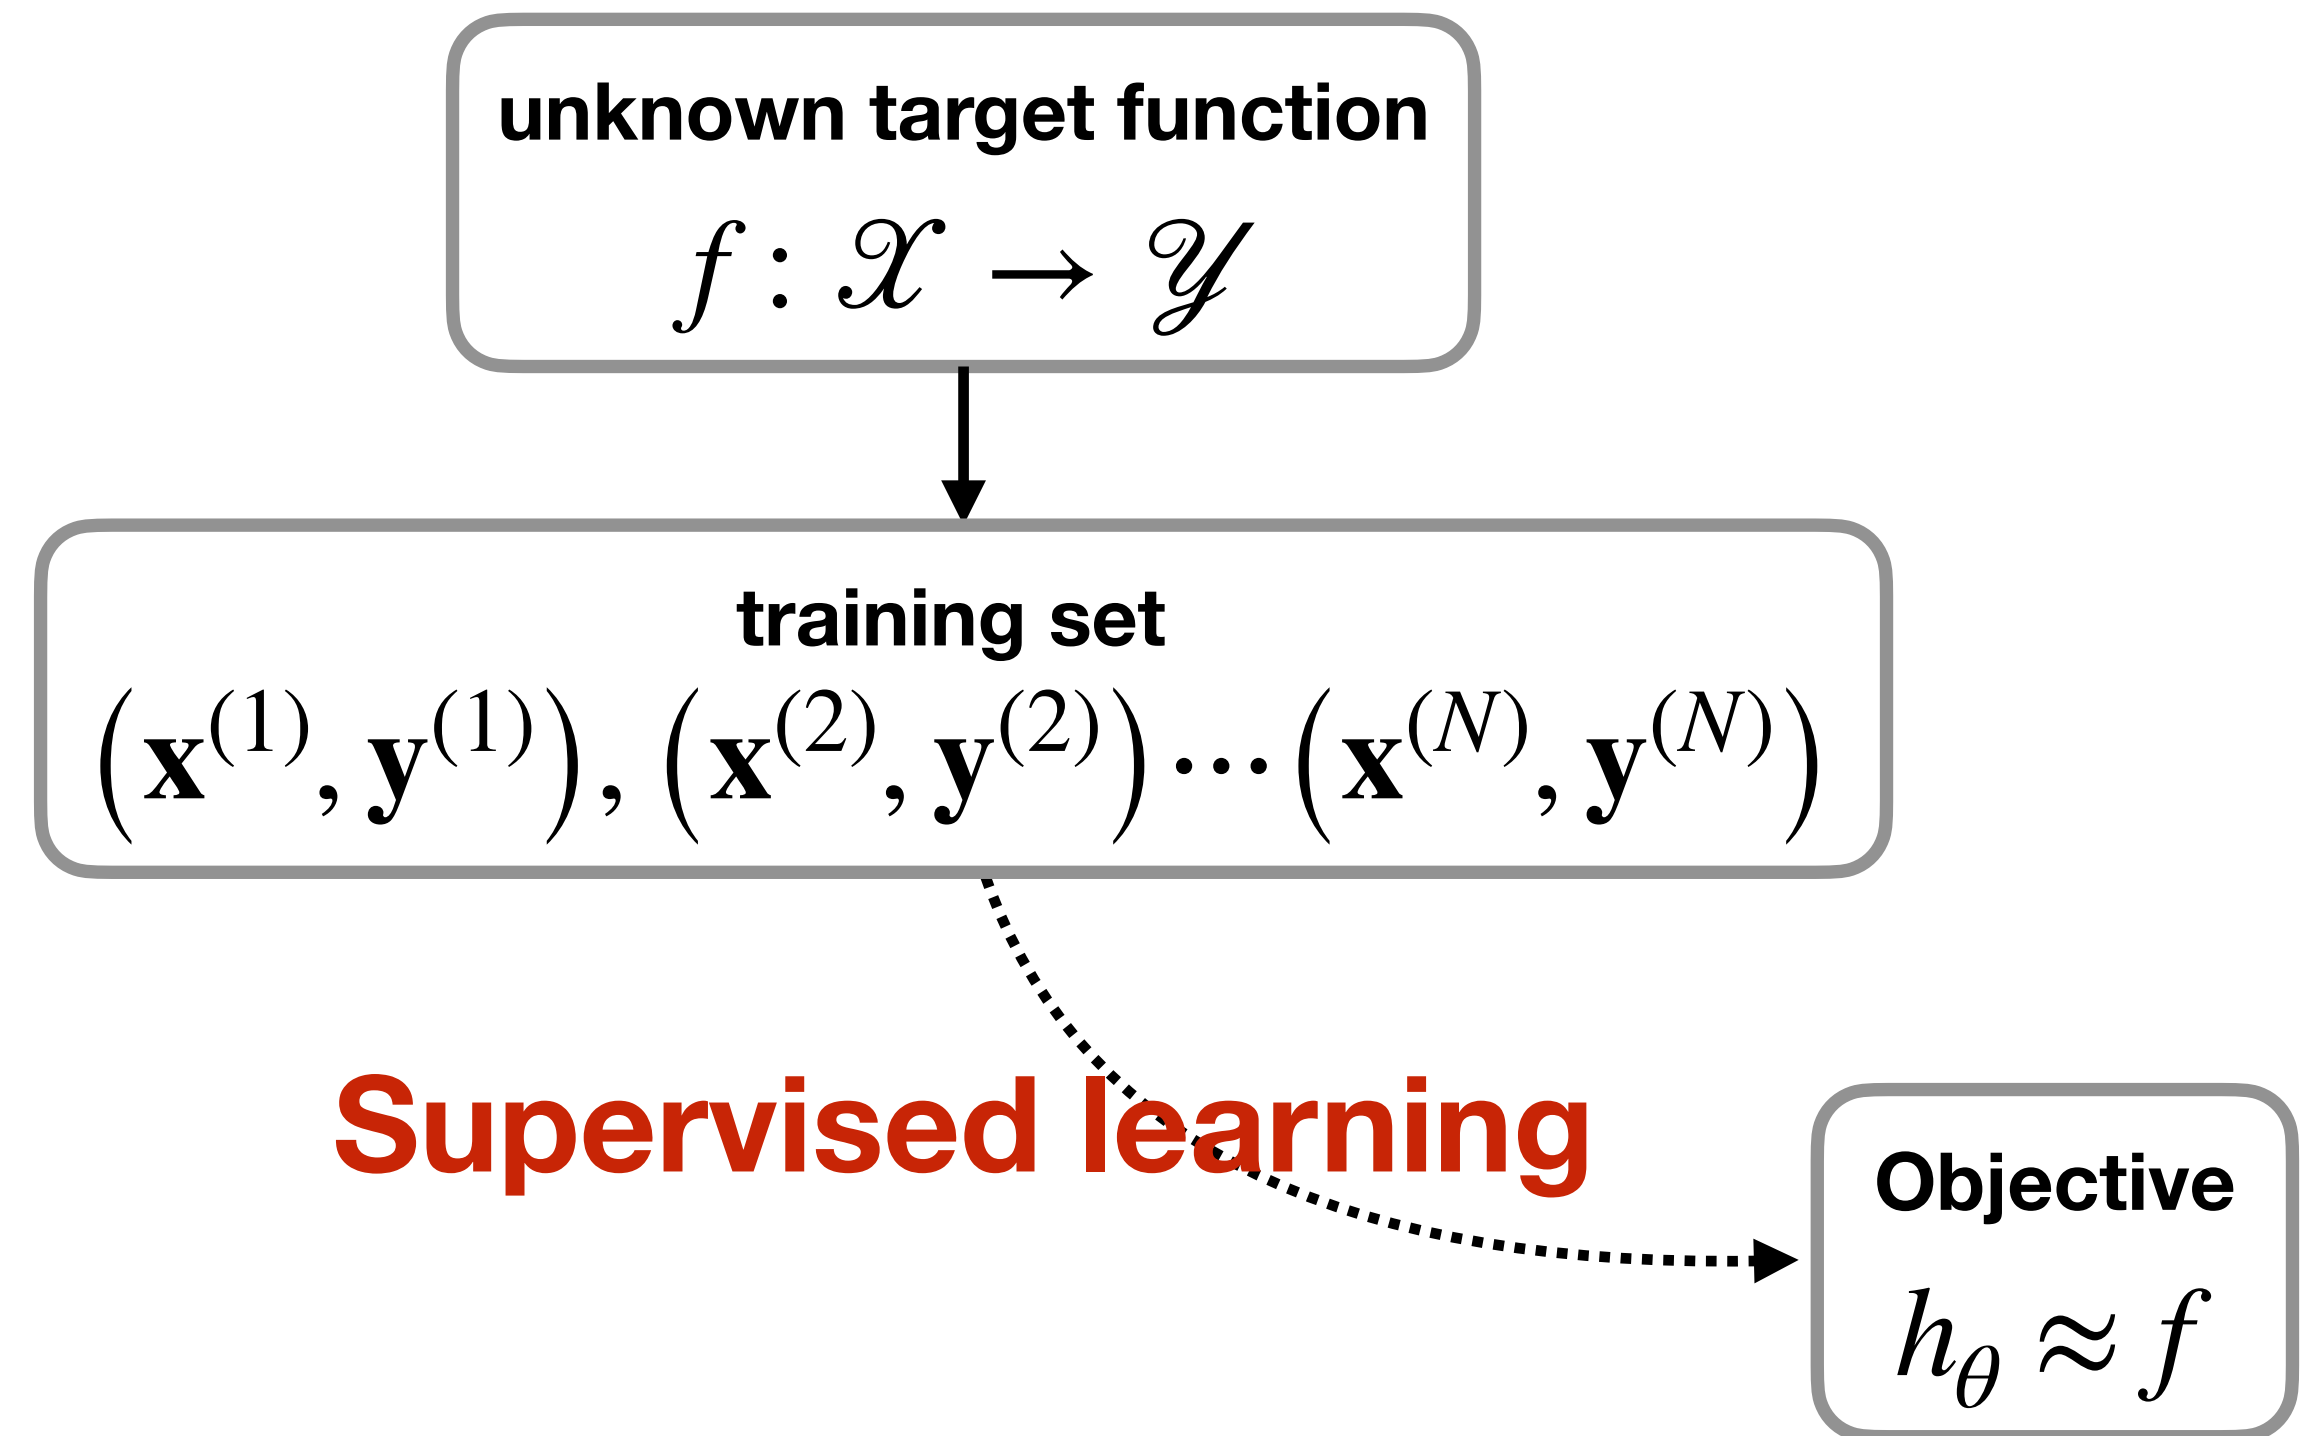





# Formalization of learning

- Input:  $\mathbf{x}$  (trunk diameter)
- Output:  $\mathbf{y}$  (leaf mass)
- Target function:  $f: \mathcal{X} \rightarrow \mathcal{Y}$  (the relationship we are looking for)
- Data:  $(\mathbf{x}^{(1)}, \mathbf{y}^{(1)}) \cdots (\mathbf{x}^{(N)}, \mathbf{y}^{(N)})$  (to be split into training and testing data)
- Hypothesis:  $h_{\theta}: \mathcal{X} \rightarrow \mathcal{Y}$  (the set of functions parametrized by  $\theta$ )

# Learning components

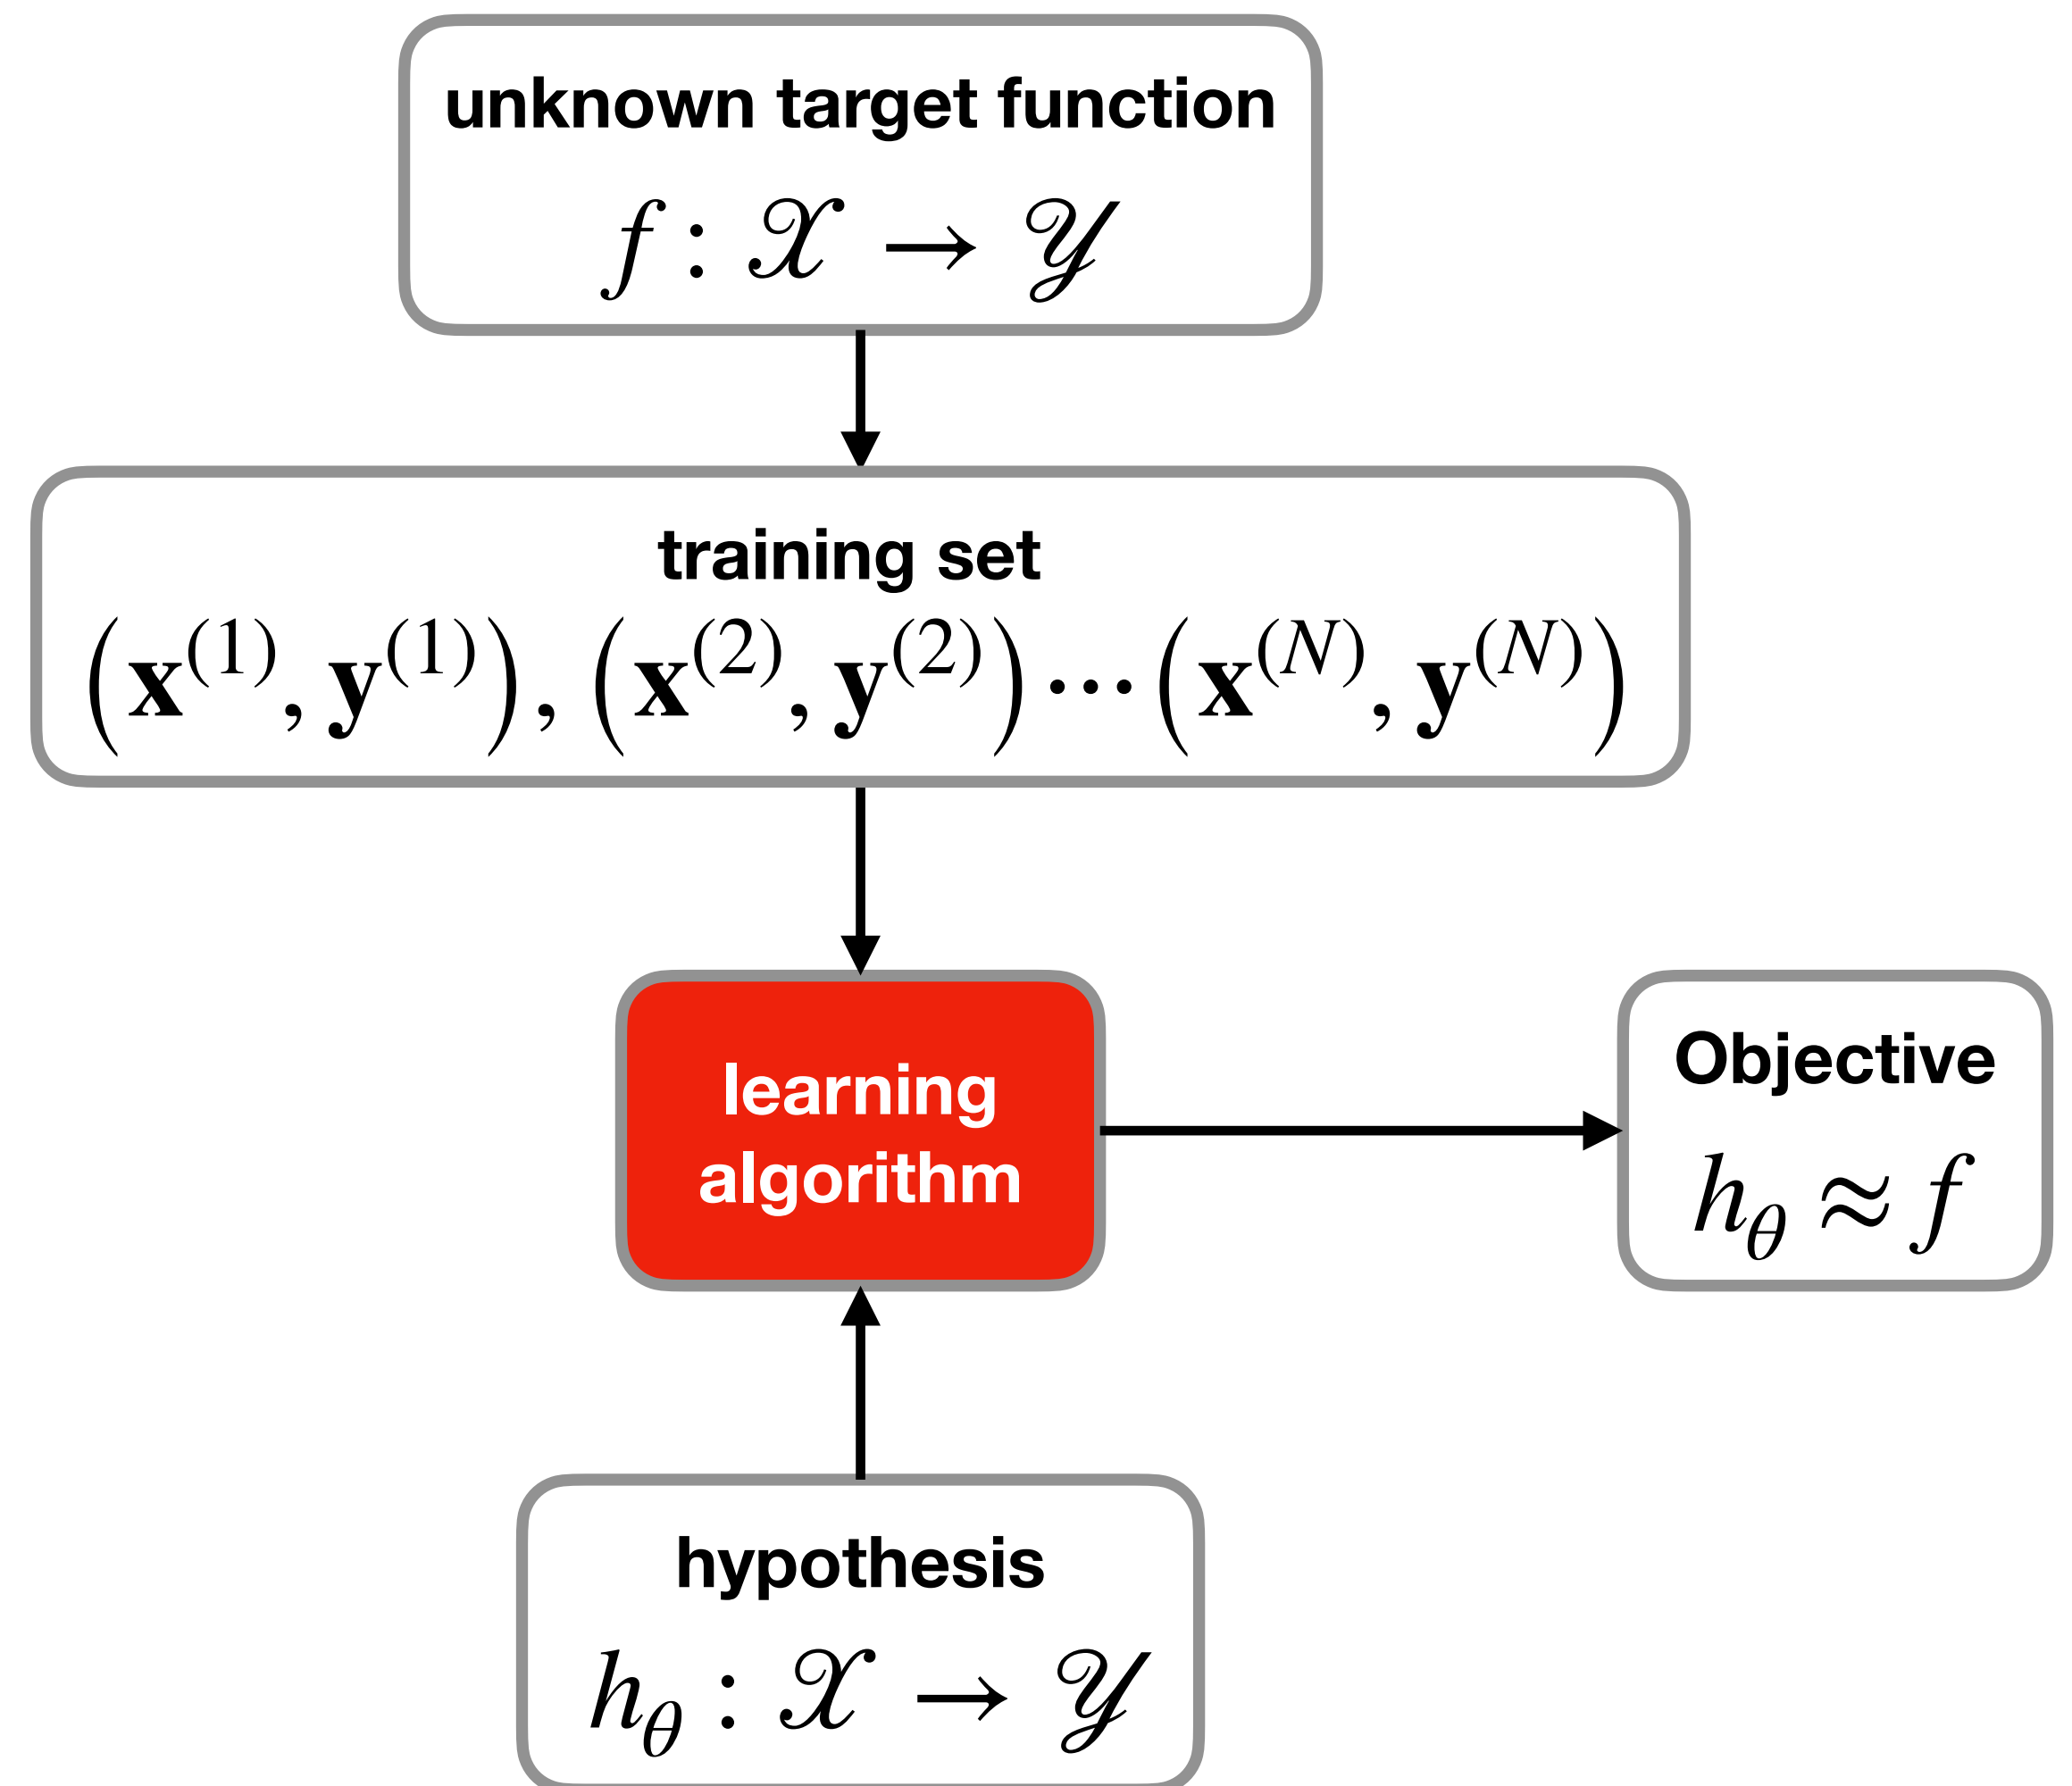


# Learning components



# Learning components

- Two components of the learning problem
  - Hypothesis:  $h_{\theta} \in \mathcal{H}$
  - Learning algorithm
    - Iterative procedure
    - Cost function
- These two components form the **learning model**



# Univariate linear regression

- Hypothesis:  $h_{\theta}(x) = \theta_0 + \theta_1 x$
- Parameters:  $\theta = [\theta_0, \theta_1]^T$
- Cost function (least squares):  $J(\theta_0, \theta_1) = \frac{1}{2N} \sum_{i=1}^N (h_{\theta}(x^{(i)}) - y^{(i)})^2$
- Goal: find  $\min_{\theta_0, \theta_1} J(\theta_0, \theta_1)$



# Gradient descent algorithm

- Iterative procedure

$$\theta_j := \theta_j - \alpha \frac{\partial}{\partial \theta_j} J(\theta)$$

$$\theta := \theta - \alpha \nabla_{\theta} J(\theta)$$

$$\frac{\partial J}{\partial \theta_0} = \frac{1}{N} \sum_{i=1}^N (h_{\theta}(x^{(i)}) - y^{(i)})$$

$$\frac{\partial J}{\partial \theta_1} = \frac{1}{N} \sum_{i=1}^N (h_{\theta}(x^{(i)}) - y^{(i)}) x^{(i)}$$

- Learning parameter:  $\alpha$
- “Batch” vs. stochastic vs. mini-batch gradient descent

# Multivariate linear regression

- Hypothesis:  $h_{\theta}(x) = \theta^T x$
- Parameters:  $\theta = [\theta_0 \cdots \theta_n]^T$
- Cost function (least squares):  $J(\theta) = \frac{1}{2N} \sum_{i=1}^N (h_{\theta}(x^{(i)}) - y^{(i)})^2$
- Goal: find  $\min_{\theta} J(\theta)$
- Gradient descent:  $\theta := \theta - \alpha \nabla_{\theta} J(\theta)$
- Normal equation:  $\theta = (X^T X)^{-1} X^T y$