

Class 10

# Deep Reinforcement Learning


Christophe Eloy

# Agent interacts with environment

At each step  $t$

- The agent:
  - Executes action  $A_t$
  - Receives observation  $O_t$
  - Receives scalar reward  $R_t$
- The environment:
  - Receives action  $A_t$
  - Emits observation  $O_{t+1}$
  - Emits scalar reward  $R_{t+1}$

Time  $t$  increments at each step



# Components of a RL agent

- **Policy** is the function that pick agent's action as a function of its state

$$a = \pi(s) \quad (\text{deterministic})$$

$$\pi(a | s) = \mathbb{P} [A_t = a | S_t = s] \quad (\text{stochastic})$$

- **Value function** is a prediction of future (discounted) rewards

$$v_{\pi}(s) = \mathbb{E}_{\pi} [R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \cdots | S_t = s]$$

- A **model** predicts what the environment will do next

$$\mathcal{P}_{ss'}^a = \mathbb{P} [S_{t+1} = s' | S_t = s, A_t = a] \quad (\text{predicts next state})$$

$$\mathcal{R}_s^a = \mathbb{E} [R_{t+1} | S_t = s, A_t = a] \quad (\text{predicts next reward})$$

# Policy and value functions

- Stochastic policy

$$\pi(a | s) = \mathbb{P} [A_t = a | S_t = s]$$

- Return:  $G_t = R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \dots$
- State-value function

$$v_{\pi}(s) = \mathbb{E}_{\pi} [G_t | S_t = s]$$

- Action-value function (Q-function)

$$q_{\pi}(s, a) = \mathbb{E}_{\pi} [G_t | S_t = s, A_t = a]$$

# Model-free prediction

- Dynamic programming
- Iterative procedure to approach state-value function, given a known policy
- Monte-Carlo algorithm (high variance, no bias)

$$V(S_t) \leftarrow V(S_t) + \alpha (G_t - V(S_t))$$

- Temporal difference algorithm (lower variance, some bias)

$$V(S_t) \leftarrow V(S_t) + \alpha (R_{t+1} + \gamma V(S_{t+1}) - V(S_t))$$

**Theorem:** MC and TD algorithm converges towards  $v_\pi(s)$

# Model-free control

- Iterative procedure to approach Q-function, given sequences  $S, A, R, S', A'$
- SARSA algorithm (on-policy)

$$Q(S, A) \leftarrow Q(S, A) + \alpha (R + \gamma Q(S', A') - Q(S, A))$$

- Q-learning algorithm (off-policy)

$$Q(S, A) \leftarrow Q(S, A) + \alpha \left( R + \gamma \max_{a'} Q(S', a') - Q(S, A) \right)$$

**Theorem:** SARSA and Q-learning converge towards  $q^*$

# DP, MC and TD

- Dynamic programming, Monte Carlo, temporal difference

- State-value function

$$V(S) \leftarrow \mathbb{E} [R + \gamma V(S') | S] \quad \text{DP (iterative policy evaluation)}$$

$$V(S) \leftarrow V(S) + \alpha (G - V(S)) \quad \text{MC}$$

$$V(S) \leftarrow V(S) + \alpha (R + \gamma V(S') - V(S)) \quad \text{TD}$$

- Q-function

$$Q(S, A) \leftarrow \mathbb{E} [R + \gamma Q(S', A') | S, A] \quad \text{DP}$$

$$Q(S, A) \leftarrow Q(S, A) + \alpha (G - Q(S, A)) \quad \text{MC}$$

$$Q(S, A) \leftarrow Q(S, A) + \alpha (R + \gamma Q(S', A') - Q(S, A)) \quad \text{TD (SARSA algorithm)}$$

# Outline

- Value function approximation
  - Implementation of a temporal difference algorithm with neural network
  - Batch methods
- Policy gradient
  - Objective functions
  - Score function
  - Policy gradient
  - Actor-critic algorithms



# Value function approximation

- We use an approximation of the action (resp. state) value function

$$\hat{q}_{\theta}(s, a) \approx q_{\pi}(s, a)$$

- Minimization of the cost:  $J(\theta) = \mathbb{E}_{\pi} \left[ \frac{1}{2} (q_{\pi}(s, a) - \hat{q}_{\theta}(s, a))^2 \right]$
- Stochastic gradient descent algorithm (SARSA)

$$\theta \leftarrow \theta + \alpha (r + \gamma \hat{q}_{\theta}(s', a') - \hat{q}_{\theta}(s, a)) \nabla_{\theta} \hat{q}_{\theta}(s, a)$$

# Policy gradient

- Objective functions

$$J_1(\theta) = \mathbb{E}_{\pi_\theta} [v(s_1)]$$

$$J_{\text{avV}}(\theta) = \mathbb{E}_{\pi_\theta} [v(s)]$$

$$J_{\text{avR}}(\theta) = \mathbb{E}_{\pi_\theta} \left[ \sum_a \pi(s, a) R_s^a \right]$$

- Score function:  $\nabla_\theta \log \pi(s, a)$

- Policy gradient

$$\theta \leftarrow \theta + \alpha \nabla_\theta J(\theta)$$

$$\nabla_\theta J(\theta) = \mathbb{E}_{\pi_\theta} \left[ \nabla_\theta \log (\pi(s, a)) Q_{\pi_\theta}(s, a) \right]$$

# Advantage actor-critic

- The critic approximates the value function (SARSA)

$$\hat{\theta} \leftarrow \hat{\theta} + \alpha \left( r + \gamma \hat{v}_{\hat{\theta}}(s') - \hat{v}_{\hat{\theta}}(s) \right) \nabla_{\hat{\theta}} \hat{v}_{\hat{\theta}}(s)$$

- The actor updates in the direction suggested by the critic (policy-gradient)

$$\theta \leftarrow \theta + \alpha \nabla_{\theta} \log \left( \pi(s, a) \right) A_{\pi_{\theta}}(s, a)$$

- Where the advantage function is

$$A_{\pi_{\theta}}(s, a) = Q_{\pi_{\theta}}(s, a) - V_{\pi_{\theta}}(s) \approx r + \gamma \hat{v}_{\hat{\theta}}(s') - \hat{v}_{\hat{\theta}}(s)$$